

# Gradient descent

From Wikipedia, the free encyclopedia

For the *analytical* method called "steepest descent", see *Method of steepest descent*.

**Gradient descent** is a **first-order iterative optimization algorithm** for finding a **local minimum** of a differentiable function. To find a local minimum of a function using gradient descent, we take steps proportional to the *negative* of the **gradient** (or approximate gradient) of the function at the current point. But if we instead take steps proportional to the *positive* of the gradient, we approach a **local maximum** of that function; the procedure is then known as **gradient ascent**. Gradient descent is generally attributed to **Cauchy**, who first suggested it in 1847,<sup>[1]</sup> but its convergence properties for non-linear optimization problems were first studied by **Haskell Curry** in 1944.<sup>[2]</sup>

## Contents [hide]

- Description
  - [An analogy for understanding gradient descent](#)
  - [Examples](#)
- [Solution of a linear system](#)
- [Solution of a non-linear system](#)
- [Comments](#)
- [Extensions](#)
  - [Fast gradient methods](#)
  - [The momentum method](#)
- [See also](#)
- [References](#)
- [Further reading](#)
- [External links](#)

## Description [\[ edit \]](#)

Gradient descent is based on the observation that if the **multi-variable function** ***F*(**x**)** is **defined** and **differentiable** in a neighborhood of a point ***a***, then ***F*(**x**)** decreases *fastest* if one goes from ***a*** in the direction of the negative gradient of ***F*** at ***a***, ***−∇F*(**a**)**. It follows that, if

$$\mathbf{a}_{n+1} = \mathbf{a}_n - \gamma \nabla F(\mathbf{a}_n)$$

for  $\gamma \in \mathbb{R}_+$  small enough, then  $F(\mathbf{a}_n) \geq F(\mathbf{a}_{n+1})$ . In other words, the term  $\gamma \nabla F(\mathbf{a})$  is subtracted from ***a*** because we want to move against the gradient, toward the local minimum. With this observation in mind, one starts with a guess ***x*<sub>0</sub>** for a local minimum of ***F***, and considers the sequence ***x*<sub>0</sub>**, ***x*<sub>1</sub>**, ***x*<sub>2</sub>**, . . . such that

$$\mathbf{x}_{n+1} = \mathbf{x}_n - \gamma_n \nabla F(\mathbf{x}_n), \; n \geq 0.$$

We have a **monotonic** sequence

$$F(\mathbf{x}_0) \geq F(\mathbf{x}_1) \geq F(\mathbf{x}_2) \geq \cdots,$$

so hopefully the sequence (***x*<sub>n</sub>**) converges to the desired local minimum. Note that the value of the *step size*  $\gamma$  is allowed to change at every iteration. With certain assumptions on the function ***F*** (for example, ***F*** **convex** and ***∇F*** **Lipschitz**) and particular choices of  $\gamma$  (e.g., chosen either via a **line search** that satisfies the **Wolfe conditions**, or the Barzilai–Borwein method<sup>[3][4]</sup> shown as following),

$$\gamma_n = \frac{|(\mathbf{x}_n - \mathbf{x}_{n-1})^T [\nabla F(\mathbf{x}_n) - \nabla F(\mathbf{x}_{n-1})]|}{\|\nabla F(\mathbf{x}_n) - \nabla F(\mathbf{x}_{n-1})\|^2}$$

**convergence** to a local minimum can be guaranteed. When the function ***F*** is **convex**, all local minima are also global minima, so in this case gradient descent can converge to the global solution.

This process is illustrated in the adjacent picture. Here ***F*** is assumed to be defined on the plane, and that its graph has a **bowl** shape. The blue curves are the **contour lines**, that is, the regions on which the value of ***F*** is constant. A red arrow originating at a point shows the direction of the negative gradient at that point. Note that the (negative) gradient at a point is **orthogonal** to the contour line going through that point. We see that gradient *descent* leads us to the bottom of the bowl, that is, to the point where the value of the function ***F*** is minimal.

## An analogy for understanding gradient descent [\[ edit \]](#)

The basic intuition behind gradient descent can be illustrated by a hypothetical scenario. A person is stuck in the mountains and is trying to get down (i.e. trying to find the global minimum). There is heavy fog such that visibility is extremely low. Therefore, the path down the mountain is not visible, so they must use local information to find the minimum. They can use the method of gradient descent, which involves looking at the steepness of the hill at their current position, then proceeding in the direction with the steepest descent (i.e. downhill). If they were trying to find the top of the mountain (i.e. the maximum), then they would proceed in the direction of steepest ascent (i.e. uphill). Using this method, they would eventually find their way down the mountain or possibly get stuck in some hole (i.e. local minimum or **saddle point**), like a mountain lake. However, assume also that the steepness of the hill is not immediately




Illustration of gradient descent on a series of level sets



Fog in the mountains

obvious with simple observation, but rather it requires a sophisticated instrument to measure, which the person happens to have at the moment. It takes quite some time to measure the steepness of the hill with the instrument, thus they should minimize their use of the instrument if they wanted to get down the mountain before sunset. The difficulty then is choosing the frequency at which they should measure the steepness of the hill so not to go off track.


In this analogy, the person represents the algorithm, and the path taken down the mountain represents the sequence of parameter settings that the algorithm will explore. The steepness of the hill represents the [slope](#) of the error surface at that point. The instrument used to measure steepness is [differentiation](#) (the slope of the error surface can be calculated by taking the [derivative](#) of the squared error function at that point). The direction they choose to travel in aligns with the [gradient](#) of the error surface at that point. The amount of time they travel before taking another measurement is the learning rate of the algorithm. See [Backpropagation § Limitations](#) for a discussion of the limitations of this type of "hill descending" algorithm.

### Examples [\[edit\]](#)

Gradient descent has problems with pathological functions such as the [Rosenbrock function](#) shown here.



$$f(x_1, x_2) = (1 - x_1)^2 + 100(x_2 - x_1^2)^2.$$

The Rosenbrock function has a narrow curved valley which contains the minimum. The bottom of the valley is very flat. Because of the curved flat valley the optimization is zigzagging slowly with small step sizes towards the minimum.



The zigzagging nature of the method is also evident below, where the gradient descent method is applied to

$$F(x, y) = \sin\left(\frac{1}{2}x^2 - \frac{1}{4}y^2 + 3\right) \cos(2x + 1 - e^y).$$



### Solution of a linear system [\[edit\]](#)

Gradient descent can be used to solve a system of linear equations, reformulated as a quadratic minimization problem, e.g., using [linear least squares](#). The solution of

$$\mathbf{Ax} - \mathbf{b} = 0$$

in the sense of linear least squares is defined as minimizing the function

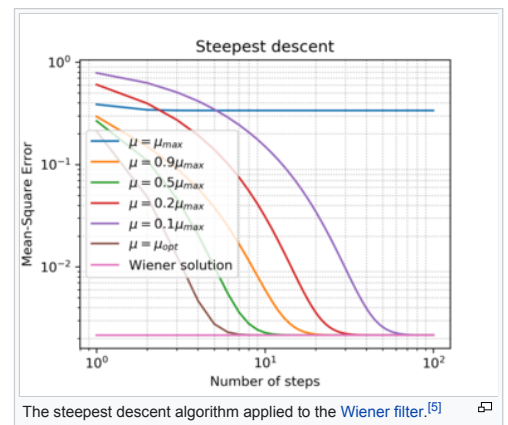
$$F(\mathbf{x}) = \|\mathbf{Ax} - \mathbf{b}\|^2.$$

In traditional linear least squares for real **A** and **b** the [Euclidean norm](#) is used, in which case

$$\nabla F(\mathbf{x}) = 2\mathbf{A}^T(\mathbf{Ax} - \mathbf{b}).$$

In this case, the [line search](#) minimization, finding the locally optimal step size **γ** on every iteration, can be performed analytically, and explicit formulas for the locally optimal **γ** are known.<sup>[6]</sup>

The algorithm is rarely used for solving linear equations, with the [conjugate gradient method](#) being one of the most popular alternatives. The number of gradient descent iterations is commonly proportional to the spectral [condition number](#)  $\kappa(\mathbf{A})$  of the system matrix **A** (the



ratio of the maximum to minimum [eigenvalues](#) of  $\mathbf{A}^T \mathbf{A}$ ), while the convergence of [conjugate gradient method](#) is typically determined by a square root of the condition number, i.e., is much faster. Both methods can benefit from [preconditioning](#), where gradient descent may require less assumptions on the preconditioner.<sup>[7]</sup>

## Solution of a non-linear system [\[ edit \]](#)

Gradient descent can also be used to solve a system of [nonlinear equations](#). Below is an example that shows how to use the gradient descent to solve for three unknown variables,  $x_1$ ,  $x_2$ , and  $x_3$ . This example shows one iteration of the gradient descent.

Consider the nonlinear system of equations

$$\begin{cases} 3x_1 - \cos(x_2 x_3) - \frac{3}{2} = 0 \\ 4x_1^2 - 625x_2^2 + 2x_2 - 1 = 0 \\ \exp(-x_1 x_2) + 20x_3 + \frac{10\pi - 3}{3} = 0 \end{cases}$$

Let us introduce the associated function

$$G(\mathbf{x}) = \begin{bmatrix} 3x_1 - \cos(x_2 x_3) - \frac{3}{2} \\ 4x_1^2 - 625x_2^2 + 2x_2 - 1 \\ \exp(-x_1 x_2) + 20x_3 + \frac{10\pi - 3}{3} \end{bmatrix},$$

where

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}.$$

One might now define the objective function

$$F(\mathbf{x}) = \frac{1}{2} G^T(\mathbf{x}) G(\mathbf{x}) = \frac{1}{2} \left[ \left( 3x_1 - \cos(x_2 x_3) - \frac{3}{2} \right)^2 + (4x_1^2 - 625x_2^2 + 2x_2 - 1)^2 + \left( \exp(-x_1 x_2) + 20x_3 + \frac{10\pi - 3}{3} \right)^2 \right],$$

which we will attempt to minimize. As an initial guess, let us use

$$\mathbf{x}^{(0)} = \mathbf{0} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}.$$

We know that

$$\mathbf{x}^{(1)} = \mathbf{0} - \gamma_0 \nabla F(\mathbf{0}) = \mathbf{0} - \gamma_0 J_G(\mathbf{0})^T G(\mathbf{0}),$$

where the [Jacobian matrix](#)  $J_G$  is given by

$$J_G(\mathbf{x}) = \begin{bmatrix} 3 & \sin(x_2 x_3) x_3 & \sin(x_2 x_3) x_2 \\ 8x_1 & -1250x_2 + 2 & 0 \\ -x_2 \exp(-x_1 x_2) & -x_1 \exp(-x_1 x_2) & 20 \end{bmatrix}.$$

We calculate:

$$J_G(\mathbf{0}) = \begin{bmatrix} 3 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 20 \end{bmatrix}, \quad G(\mathbf{0}) = \begin{bmatrix} -2.5 \\ -1 \\ 10.472 \end{bmatrix}.$$

Thus

$$\mathbf{x}^{(1)} = \mathbf{0} - \gamma_0 \begin{bmatrix} -7.5 \\ -2 \\ 209.44 \end{bmatrix},$$

and

$$F(\mathbf{0}) = 0.5 \left( (-2.5)^2 + (-1)^2 + (10.472)^2 \right) = 58.456.$$

Now, a suitable  $\gamma_0$  must be found such that

$$F(\mathbf{x}^{(1)}) \leq F(\mathbf{x}^{(0)}) = F(\mathbf{0}).$$

This can be done with any of a variety of [line search](#) algorithms. One might also simply guess  $\gamma_0 = 0.001$ , which gives


$$\mathbf{x}^{(1)} = \begin{bmatrix} 0.0075 \\ 0.002 \\ -0.20944 \end{bmatrix}.$$

Evaluating the objective function at this value, yields

$$F(\mathbf{x}^{(1)}) = 0.5 \left( (-2.48)^2 + (-1.00)^2 + (6.28)^2 \right) = 23.306.$$

The decrease from  $F(\mathbf{0}) = 58.456$  to the next step's value of

$$F(\mathbf{x}^{(1)}) = 23.306$$



An animation showing the first 83 iterations of gradient descent applied to this example. Surfaces are [isosurfaces](#) of  $F(\mathbf{x}^{(n)})$  at current guess  $\mathbf{x}^{(n)}$ , and arrows show the direction of descent. Due to a small and constant step size, the convergence is slow.

is a sizable decrease in the objective function. Further steps would reduce its value further, until an approximate solution to the system was found.

## Comments [[edit](#)]

Gradient descent works in spaces of any number of dimensions, even in infinite-dimensional ones. In the latter case, the search space is typically a [function space](#), and one calculates the [Fréchet derivative](#) of the functional to be minimized to determine the descent direction.<sup>[8]</sup>

That gradient descent works in any number of dimensions (finite number at least) can be seen as a consequence of the [Cauchy-Schwarz inequality](#). That article proves that the magnitude of the inner (dot) product of two vectors of any dimension is maximized when they are colinear. In the case of gradient descent, that would be when the vector of independent variable adjustments is proportional to the gradient vector of partial derivatives.

The gradient descent can take many iterations to compute a local minimum with a required [accuracy](#), if the [curvature](#) in different directions is very different for the given function. For such functions, [preconditioning](#), which changes the geometry of the space to shape the function level sets like [concentric circles](#), cures the slow convergence. Constructing and applying preconditioning can be computationally expensive, however.

The gradient descent can be combined with a [line search](#), finding the locally optimal step size  $\gamma$  on every iteration. Performing the line search can be time-consuming. Conversely, using a fixed small  $\gamma$  can yield poor convergence.

Methods based on [Newton's method](#) and inversion of the [Hessian](#) using [conjugate gradient](#) techniques can be better alternatives.<sup>[9][10]</sup> Generally, such methods converge in fewer iterations, but the cost of each iteration is higher. An example is the [BFGS method](#) which consists in calculating on every step a matrix by which the gradient vector is multiplied to go into a "better" direction, combined with a more sophisticated [line search](#) algorithm, to find the "best" value of  $\gamma$ . For extremely large problems, where the computer-memory issues dominate, a limited-memory method such as [L-BFGS](#) should be used instead of BFGS or the steepest descent.

Gradient descent can be viewed as applying [Euler's method](#) for solving [ordinary differential equations](#)  $\mathbf{x}'(t) = -\nabla f(\mathbf{x}(t))$  to a [gradient flow](#). In turn, this equation may be derived as an optimal controller<sup>[11]</sup> for the control system  $\mathbf{x}'(t) = \mathbf{u}(t)$  with  $\mathbf{u}(t)$  given in feedback form  $\mathbf{u}(t) = -\nabla f(\mathbf{x}(t))$ .

## Extensions [[edit](#)]

Gradient descent can be extended to handle [constraints](#) by including a [projection](#) onto the set of constraints. This method is only feasible when the projection is efficiently computable on a computer. Under suitable assumptions, this method converges. This method is a specific case of the [forward-backward algorithm](#) for monotone inclusions (which includes [convex programming](#) and [variational inequalities](#)).<sup>[12]</sup>

### Fast gradient methods [[edit](#)]

Another extension of gradient descent is due to [Yurii Nesterov](#) from 1983,<sup>[13]</sup> and has been subsequently generalized. He provides a simple modification of the algorithm that enables faster convergence for convex problems. For unconstrained smooth problems the method is called the [Fast Gradient Method](#) (FGM) or the [Accelerated Gradient Method](#) (AGM). Specifically, if the differentiable function  $F$  is convex and  $\nabla F$  is [Lipschitz](#), and it is not assumed that  $F$  is [strongly convex](#), then the error in the objective value generated at each step  $k$  by the gradient descent method will be [bounded by](#)  $\mathcal{O}\left(\frac{1}{k}\right)$ . Using the Nesterov acceleration technique, the error decreases at  $\mathcal{O}\left(\frac{1}{k^2}\right)$ .<sup>[14]</sup> It is known that the rate  $\mathcal{O}(k^{-2})$  for the decrease of the [cost function](#) is optimal for first-order optimization methods. Nevertheless, there is the opportunity to improve the algorithm by reducing the constant factor. The [optimized gradient method](#) (OGM)<sup>[15]</sup> reduces that constant by a factor of two and is an optimal first-order method for large-scale problems.<sup>[16]</sup>

For constrained or non-smooth problems, Nesterov's FGM is called the [fast proximal gradient method](#) (FPGM), an acceleration of the [proximal gradient method](#).

### The momentum method [[edit](#)]

*Main article: [Stochastic gradient descent § Momentum](#)*

Yet another extension, that reduces the risk of getting stuck in a local minimum, as well as speeds up the convergence considerably in cases where the process would otherwise zig-zag heavily, is the *momentum method*, which uses a momentum term in analogy to "the mass of Newtonian particles that move through a viscous medium in a conservative force field".<sup>[17]</sup> This method is often used as an extension to the [backpropagation](#) algorithms used to train [artificial neural networks](#).<sup>[18][19]</sup>

## See also [[edit](#)]

- |  |   |   |
|--|---|---|
| <ul style="list-style-type: none"><li><a href="#">Conjugate gradient method</a></li><li><a href="#">Stochastic gradient descent</a></li><li><a href="#">Rprop</a></li><li><a href="#">Delta rule</a></li></ul> | <ul style="list-style-type: none"><li><a href="#">Wolfe conditions</a></li><li><a href="#">Preconditioning</a></li><li><a href="#">Broyden–Fletcher–Goldfarb–Shanno algorithm</a></li><li><a href="#">Davidon–Fletcher–Powell formula</a></li></ul> | <ul style="list-style-type: none"><li><a href="#">Nelder–Mead method</a></li><li><a href="#">Gauss–Newton algorithm</a></li><li><a href="#">Hill climbing</a></li><li><a href="#">Quantum annealing</a></li></ul> |
|--|---|---|

## References [[edit](#)]

1. <span><span>↑</span></span> <span>Lemaréchal, C. (2012). "Cauchy and the Gradient Method" <span><span><span></span></span></span> (PDF). <i>Doc Math Extra</i>: 251–254.</span>	6. <span><span>↑</span></span> <span>Yuan, Ya-xiang (1999). "Step-sizes for the gradient method" <span><span><span></span></span></span> (PDF). <i>AMS/IP Studies in Advanced Mathematics</i>. <b>42</b> (2): 785.</span>
2. <span><span>↑</span></span> <span>Curry, Haskell B. (1944). "The Method of Steepest Descent for Non-linear Minimization Problems". <i>Quart. Appl. Math.</i> <b>2</b> (3): 258–261. doi:10.1090/qam/10667<span><span><span></span></span></span>.</span>	7. <span><span>↑</span></span> <span>Henricus Bouwmeester, Andrew Dougherty, Andrew V Knyazev. Nonsymmetric Preconditioning for Conjugate Gradient and Steepest Descent Methods. <i>Procedia Computer Science</i>, Volume 51, Pages 276-285, Elsevier, 2015. <a href="https://doi.org/10.1016/j.procs.2015.05.241">https://doi.org/10.1016/j.procs.2015.05.241</a><span><span><span></span></span></span></span>
3. <span><span>↑</span></span> <span>Barzilai, Jonathan; Borwein, Jonathan M. (1988). "Two-Point Step Size Gradient Methods". <i>IMA Journal of Numerical Analysis</i>. <b>8</b> (1): 141–148. doi:10.1093/imanum/8.1.141<span><span><span></span></span></span>.</span>	8. <span><span>↑</span></span> <span>Akilov, G. P.; Kantorovich, L. V. (1982). <i>Functional Analysis</i> (2nd ed.). Pergamon Press. ISBN<span> </span>0-08-023036-9.</span>
4. <span><span>↑</span></span> <span>Fletcher, R. (2005). "On the Barzilai–Borwein Method". In Qi, L.; Teo, K.; Yang, X. (eds.). <i>Optimization and Control with Applications</i>. Applied Optimization. <b>96</b>. Boston: Springer. pp.<span> </span>235–256. ISBN<span> </span>0-387-24254-6.</span>	9. <span><span>↑</span></span> <span>Press, W. H.; Teukolsky, S. A.; Vetterling, W. T.; Flannery, B. P. (1992). <i>Numerical Recipes in C: The Art of Scientific Computing</i> <span><span><span></span></span></span> (2nd ed.). New York: Cambridge University Press. ISBN<span> </span>0-521-43108-5.</span>
5. <span><span>↑</span></span> <span>Haykin, Simon S. Adaptive filter theory. Pearson Education India, 2008. - p. 108-142, 217-242</span>	10. <span><span>↑</span></span> <span>Strutz, T. (2016). <i>Data Fitting and Uncertainty: A Practical Introduction to Weighted Least Squares and Beyond</i> (2nd ed.). Springer Vieweg. ISBN<span> </span>978-3-658-11455-8.</span>

11. <sup>^</sup> Ross, I. M. (2019-07-01). "An optimal control theory for nonlinear optimization"<sup>↗</sup>. *Journal of Computational and Applied Mathematics*. **354**: 39–51. doi:10.1016/j.cam.2018.12.044<sup>↗</sup>. ISSN 0377-0427<sup>↗</sup>.

12. <sup>^</sup> Combettes, P. L.; Pesquet, J.-C. (2011). "Proximal splitting methods in signal processing". In Bauschke, H. H.; Burachik, R. S.; Combettes, P. L.; Elser, V.; Luke, D. R.; Wolkowicz, H. (eds.). *Fixed-Point Algorithms for Inverse Problems in Science and Engineering*. New York: Springer. pp. 185–212. arXiv:0912.3522<sup>↗</sup>. ISBN 978-1-4419-9568-1.

13. <sup>^</sup> Nesterov, Yu. (2004). *Introductory Lectures on Convex Optimization : A Basic Course*. Springer. ISBN 1-4020-7553-7.

14. <sup>^</sup> Vandenberghe, Lieven (2019). "Fast Gradient Methods"<sup>↗</sup> (PDF). *Lecture notes for EE236C at UCLA*.

15. <sup>^</sup> Kim, D.; Fessler, J. A. (2016). "Optimized First-order Methods for Smooth Convex Minimization". *Math. Prog.* **151** (1–2): 81–107. arXiv:1406.5468<sup>↗</sup>. doi:10.1007/s10107-015-0949-3<sup>↗</sup>.

16. <sup>^</sup> Drori, Yoel (2017). "The Exact Information-based Complexity of Smooth Convex Minimization". *Journal of Complexity*. **39**: 1–16. arXiv:1606.01424<sup>↗</sup>. doi:10.1016/j.jco.2016.11.001<sup>↗</sup>.

17. <sup>^</sup> Qian, Ning (January 1999). "On the momentum term in gradient descent learning algorithms"<sup>↗</sup> (PDF). *Neural Networks*. **12** (1): 145–151. CiteSeerX 10.1.1.57.5612<sup>↗</sup>. doi:10.1016/S0893-6080(98)00116-6<sup>↗</sup>. Archived from the original<sup>↗</sup> (PDF) on 8 May 2014. Retrieved 17 October 2014.

18. <sup>^</sup> "Momentum and Learning Rate Adaptation"<sup>↗</sup>. Willamette University. Retrieved 17 October 2014.

19. <sup>^</sup> Geoffrey Hinton; Nitish Srivastava; Kevin Swersky. "The momentum method"<sup>↗</sup>. *Coursera*. Retrieved 2 October 2018. Part of a lecture series for the Coursera online course *Neural Networks for Machine Learning*<sup>↗</sup> Archived<sup>↗</sup> 2016-12-31 at the Wayback Machine.

Further reading <sup>[ edit ]</sup>

• Boyd, Stephen; Vandenberghe, Lieven (2004). "Unconstrained Minimization"<sup>↗</sup> (PDF). *Convex Optimization*. New York: Cambridge University Press. pp. 457–520. ISBN 0-521-83378-7.

External links <sup>[ edit ]</sup>

• Using gradient descent in C++, Boost, Ublas for linear regression<sup>↗</sup>

• Series of Khan Academy videos discusses gradient ascent<sup>↗</sup>

• Online book teaching gradient descent in deep neural network context<sup>↗</sup>

• "Gradient Descent, How Neural Networks Learn"<sup>↗</sup>. *3Blue1Brown*. October 16, 2017 – via YouTube.

<span>V</span> <span>•</span> <span>T</span> <span>•</span> <span>E</span>	<b>Optimization: Algorithms, methods, and heuristics</b>	<span>[</span> hide <span>]</span>
	<b>Unconstrained nonlinear</b>	<span>[</span> show <span>]</span>
	<b>Constrained nonlinear</b>	<span>[</span> show <span>]</span>
	<b>Convex optimization</b>	<span>[</span> show <span>]</span>
	<b>Combinatorial</b>	<span>[</span> show <span>]</span>
	<b>Metaheuristics</b>	<span>[</span> show <span>]</span>
	Software	
<span>V</span> <span>•</span> <span>T</span> <span>•</span> <span>E</span>	<b>Differentiable computing</b>	<span>[</span> show <span>]</span>

Categories: Mathematical optimization | First order methods | Optimization algorithms and methods | Gradient methods

This page was last edited on 2 August 2020, at 12:50 (UTC).

Text is available under the Creative Commons Attribution-ShareAlike License; additional terms may apply. By using this site, you agree to the Terms of Use and Privacy Policy. Wikipedia® is a registered trademark of the Wikimedia Foundation, Inc., a non-profit organization.

Privacy policy

About Wikipedia

Disclaimers

Contact Wikipedia

Mobile view

Developers

Statistics

Cookie statement